

optimisation in neural networks

Patrick van der Smagt

What is optimisation?

We have a model $p_{\mathbf{w}}(z \mid x)$. This can, e.g., be a neural network.



We want to minimise the loss

$$\mathcal{L}(\mathbf{w}) = -\log \prod_i p_{\mathbf{w}}(z_i \mid x_i)$$

by finding better values of \mathbf{w} .

How do we do optimisation?

- if finding the best \mathbf{w} is a convex problem, good methods exist (remember SVD from linear algebra).
- In general, finding the best \mathbf{w} is not a convex problem. Only incremental methods are known.



Convex optimisation problems

Practical example: $\{(x_i, z_i)\} = \{(0, 1); (1, 2.1); (2, 2.9)\}$.

Our model: $y_{\mathbf{w}}(x) = ax + b$ with $\mathbf{w} = (a, b)$; the MLE loss is


$$\mathcal{L}(\mathbf{w}) = \sum_i (y_{\mathbf{w}}(x_i) - z_i)^2$$

How do we find the minimum of \mathcal{L} ? It is there where $\partial\mathcal{L}/\partial w_i = 0$.

$$\frac{\partial\mathcal{L}_i(\mathbf{w})}{\partial w_1 \equiv a} = 2x_i(b + ax_i - z_i) = 0$$


$$\frac{\partial\mathcal{L}_i(\mathbf{w})}{\partial w_2 \equiv b} = 2(b + ax_i - z_i) = 0$$

We can solve that!



What if...

$$y_{(a,b,c)}(x) = a \exp(-b(\mathbf{x} - c)^2) + d$$




We can't find a closed-form solution for that!

$$2e^{-bc^2} \left(ae^{-bc^2} - 1 \right) + 2e^{-b(1-c)^2} \left(ae^{-b(1-c)^2} - 2.1 \right) + \\ 2e^{-b(2-c)^2} \left(ae^{-b(2-c)^2} - 2.9 \right) = 0$$

$$-2ac^2e^{-bc^2} \left(ae^{-bc^2} - 1 \right) - 2a(1-c)^2e^{-b(1-c)^2} \left(ae^{-b(1-c)^2} - 2.1 \right) - \\ 2a(2-c)^2e^{-b(2-c)^2} \left(ae^{-b(2-c)^2} - 2.9 \right) = 0$$

$$-4abce^{-bc^2} \left(ae^{-bc^2} - 1 \right) + 4ab(1-c)e^{-b(1-c)^2} \left(ae^{-b(1-c)^2} - 2.1 \right) + \\ 4ab(2-c)e^{-b(2-c)^2} \left(ae^{-b(2-c)^2} - 2.9 \right) = 0$$




But

We can compute $\partial \mathcal{L}(\mathbf{w}) / \partial w_i$

the value of \mathcal{L}


we are interested in finding $\arg \min_{\mathbf{w}} \mathcal{L}$



using local information $\mathcal{L}(\mathbf{w})$

we are interested in finding $\arg \min_{\mathbf{w}} \mathcal{L}$


we usually only have local information



using local information $\mathcal{L}(\mathbf{w})$ as well as $\partial\mathcal{L}(\mathbf{w})/\partial w$

we are interested in finding $\arg \min_{\mathbf{w}} \mathcal{L}$


we usually only have local information



using the gradient \mathbf{g} of \mathcal{L}

The direction \mathbf{u} in which to optimise is given by the gradient:

$$\mathbf{u} = -\mathbf{g}$$



Searching the minimum by repeated evaluation of \mathcal{L} and $\mathbf{g} \equiv \nabla \mathcal{L}$.

$-\mathbf{g}$ gives us a direction \mathbf{u} in which we want to optimise.

We change the parameter vector as follows:

$$\mathbf{u}_i = -\mathbf{g}_i \tag{1}$$

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \alpha \mathbf{u}_i \tag{2}$$

we call \mathbf{u} the **search direction**

we call α the **learning parameter** or **step size**


we call this method **steepest descent** or **gradient descent**

it belongs to the class of greedy algorithms

the value of α

a *too small* value for α has two drawbacks:


- we find the minimum more slowly
- we end up in local minima or saddle/flat points



the value of α


a *too large* value for α has one drawback:

- you may never find a minimum; oscillations usually occur



we only need 2 steps to overshoot!

putting a trace on \mathbf{u}



a: $\mathbf{u} = -\mathbf{g}$; small α ;

b: $\mathbf{u} = -\mathbf{g}$; large α ;

c:

$$\mathbf{u}_0 = -\mathbf{g}_0 \quad (3)$$

$$\mathbf{u}_i = -\mathbf{g}_i + \beta \mathbf{u}_{i-1} \quad (4)$$

$$= -\mathbf{g}_i - \beta \mathbf{g}_{i-1} - \beta^2 \mathbf{g}_{i-2} - \beta^3 \mathbf{g}_{i-3} \dots \quad (5)$$

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \alpha \mathbf{u}_i \quad (6)$$

we call α the **learning rate**

we call β the **momentum**

we usually take $\beta \gg \alpha$

trick: momentum

How do we choose α and β ? if, for the sake of the argument, assume that $\mathbf{g} \equiv \nabla E$ does not change:


$$\begin{aligned}\Delta \mathbf{w} &= -\alpha \mathbf{g} (1 + \beta + \beta^2 + \dots) \\ &= -\frac{\alpha}{1 - \beta} \mathbf{g}\end{aligned}$$

Assuming a perfect ∇E , the best values for α and β are when


$$\frac{\alpha}{1 - \beta} = 1 \quad \Rightarrow \quad \alpha + \beta = 1$$

Typically we choose α small and β large (of course, $\alpha, \beta > 0$).

bird's eye view




=



$$\mathcal{L}(x) = \mathcal{L}(0) + x \underbrace{\frac{\partial \mathcal{L}}{\partial w}}_g + x^2 \underbrace{\frac{\partial^2 \mathcal{L}}{\partial w^2}}_{\text{Hessian} H}$$

optimising

following the gradient is not always the best choice




Close to minima, it appears that Loss functions are close to quadratic


condition of the Hessian

Condition = largest EV / smallest EV

Condition 5



Condition 100



What does H look like?

A large condition number means that some directions of H are very steep compared to others. In neural networks, a condition of 10^{10} is not uncommon.

A class of optimisers (CG, Adam, rprop, adadelata, ...) deal with such H .

One resource I like in particular:

<https://www.benfrederickson.com/numerical-optimization/>